

Comparitive Performance Report for directreadsexperiment vs directreadsexperiment

Table of contents

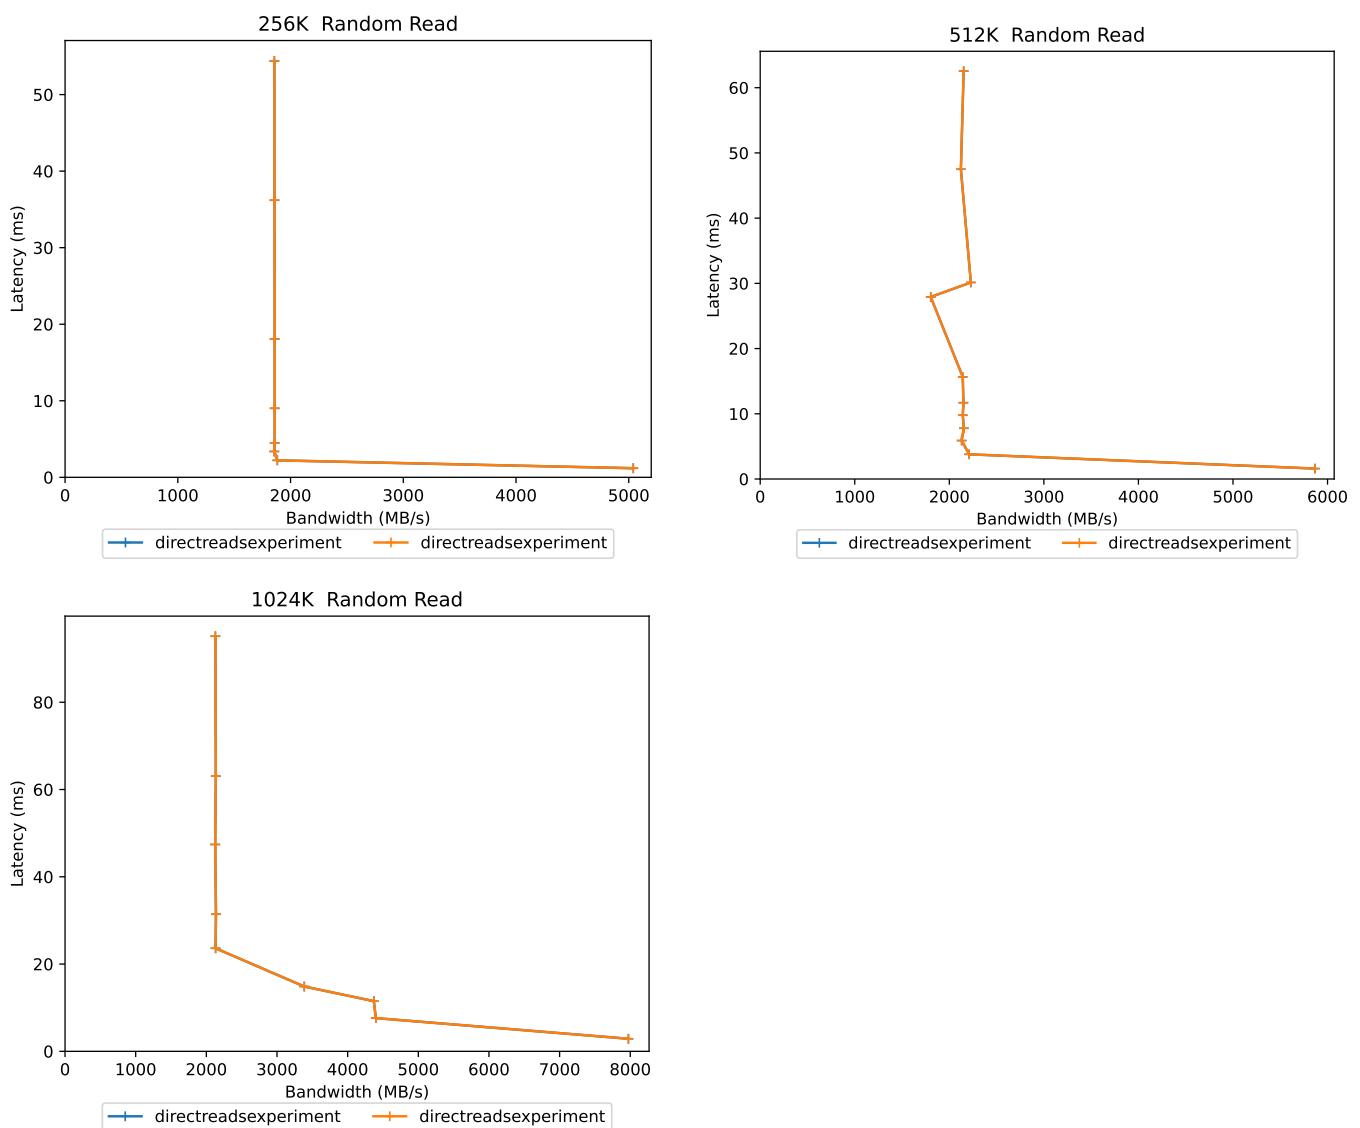
- Comparison summary for directreadsexperiment vs directreadsexperiment
- Response Curves
 - Random Read
- Configuration yaml files
 - results

Comparison summary for directreadsexperiment vs directreadsexperiment

Random Read	directreadsexperiment	directreadsexperiment	%change throughput	%change latency
256K	5040 MB/s@1.2ms	5040 MB/s@1.2ms	0%	0%
512K	5867 MB/s@1.6ms	5867 MB/s@1.6ms	0%	0%
1024K	7974 MB/s@2.9ms	7974 MB/s@2.9ms	0%	0%

Response Curves

Random Read



Configuration yaml files

Only yaml files that differ by more than 20 lines from the yaml file for the baseline directory will be added here in addition to the baseline yaml

results

```
librbd_fio:  
  cmd_path: /usr/local/bin/fio2  
  create_report: true  
  fio_out_format: json  
  log_avg_msec: 100  
  log_bw: true  
  log_iops: true  
  log_lat: true  
  norandommap: true  
  osd_ra:  
    - 4096  
  poolname: rbd_replicated  
  prefill:  
    blocksize: 64k  
    numjobs: 1  
  procs_per_volume:  
    - 1  
  ramp: 20  
  rbdname: cbt-librbd_fio  
  time: 60  
  time_based: true  
  use_existing_volumes: true  
  vol_size: 1000  
  volumes_per_client:  
    - 8  
  wait_pgautoscaler_timeout: 10  
workloads:  
  1Mrandomread:  
    jobname: randread  
    mode: randread  
    numjobs:  
      - 1  
    op_size: 1048576  
    total_iodepth:  
      - 1  
      - 2  
      - 3  
      - 4  
      - 8  
      - 12  
      - 16  
      - 20  
      - 24  
      - 28  
      - 32  
      - 48  
      - 64  
      - 96  
      - 128  
      - 192  
  256krandomread:  
    jobname: randread  
    mode: randread  
    numjobs:  
      - 1  
    op_size: 262144  
    total_iodepth:  
      - 1  
      - 2  
      - 3  
      - 4
```

```
- 5
- 8
- 16
- 24
- 32
- 64
- 128
- 256
- 384
512krandomread:
  jobname: randread
  mode: randread
  numjobs:
    - 1
  op_size: 524288
  total_iodepth:
    - 1
    - 2
    - 3
    - 4
    - 8
    - 16
    - 24
    - 32
    - 40
    - 48
    - 64
    - 96
    - 128
    - 192
    - 256
cluster:
  archive_dir: /tmp/cbt
  ceph-mgr_cmd: /usr/bin/ceph-mgr
  ceph-mon_cmd: /usr/bin/ceph-mon
  ceph-osd_cmd: /usr/bin/ceph-osd
  ceph-run_cmd: /usr/bin/ceph-run
  ceph_cmd: /usr/bin/ceph
  clients:
    - --- server1 ---
  clusterid: ceph
  conf_file: /cbt/ceph.conf.4x1x1.fs
  fs: xfs
  head: --- server1 ---
  iterations: 1
  mgrs:
    --- server1 ---:
      a: null
  mkfs_opts: -f -i size=2048
  mons:
    --- server1 ---:
      a: --- IP Address --:6789
  mount_opts: -o inode64,noatime,logbsize=256k
  osds:
    - --- server1 ---
    osds_per_node: 6
    pdsh_ssh_args: -a -x -l%u %h
    rados_cmd: /usr/bin/rados
    rbd_cmd: /usr/bin/rbd
    tmp_dir: /tmp/cbt
    use_existing: true
    user: root
  monitoring_profiles:
```

```
collectl:
  args: -c 18 -sCD -i 10 -P -oz -F0 --rawtoo --sep ";" -f {collectl_dir}

librbd_fio:
  cmd_path: /usr/local/bin/fio2
  create_report: true
  fio_out_format: json
  log_avg_msec: 100
  log_bw: true
  log_iops: true
  log_lat: true
  norandommap: true
  osd_ra:
    - 4096
  poolname: rbd_replicated
  prefill:
    blocksize: 64k
    numjobs: 1
  procs_per_volume:
    - 1
  ramp: 20
  rbdname: cbt-librbd_fio
  time: 60
  time_based: true
  use_existing_volumes: true
  vol_size: 1000
  volumes_per_client:
    - 8
  wait_pgautoscaler_timeout: 10
workloads:
  1Mrandomread:
    jobname: randread
    mode: randread
    numjobs:
      - 1
    op_size: 1048576
    total_iodepth:
      - 1
      - 2
      - 3
      - 4
      - 8
      - 12
      - 16
      - 20
      - 24
      - 28
      - 32
      - 48
      - 64
      - 96
      - 128
      - 192
  256krandomread:
    jobname: randread
    mode: randread
    numjobs:
      - 1
    op_size: 262144
    total_iodepth:
      - 1
      - 2
      - 3
```

```
- 4
- 5
- 8
- 16
- 24
- 32
- 64
- 128
- 256
- 384
512krandomread:
  jobname: randread
  mode: randread
  numjobs:
    - 1
  op_size: 524288
  total_iodepth:
    - 1
    - 2
    - 3
    - 4
    - 8
    - 16
    - 24
    - 32
    - 40
    - 48
    - 64
    - 96
    - 128
    - 192
    - 256
cluster:
  archive_dir: /tmp/cbt
  ceph-mgr_cmd: /usr/bin/ceph-mgr
  ceph-mon_cmd: /usr/bin/ceph-mon
  ceph-osd_cmd: /usr/bin/ceph-osd
  ceph-run_cmd: /usr/bin/ceph-run
  ceph_cmd: /usr/bin/ceph
  clients:
    - --- server1 ---
  clusterid: ceph
  conf_file: /cbt/ceph.conf.4x1x1.fs
  fs: xfs
  head: --- server1 ---
  iterations: 1
  mgrs:
    --- server1 ---:
      a: null
  mkfs_opts: -f -i size=2048
  mons:
    --- server1 ---:
      a: --- IP Address --:6789
  mount_opts: -o inode64,noatime,logbsize=256k
  osds:
    - --- server1 ---
  osds_per_node: 6
  pdsh_ssh_args: -a -x -l%u %h
  rados_cmd: /usr/bin/rados
  rbd_cmd: /usr/bin/rbd
  tmp_dir: /tmp/cbt
  use_existing: true
  user: root
```

```
monitoring_profiles:  
  collectl:  
    args: -c 18 -sCD -i 10 -P -oz -F0 --rawtoo --sep ";" -f {collectl_dir}
```