

Comparitive Performance Report for directreadsexperiment vs directreadsexperiment

Table of contents

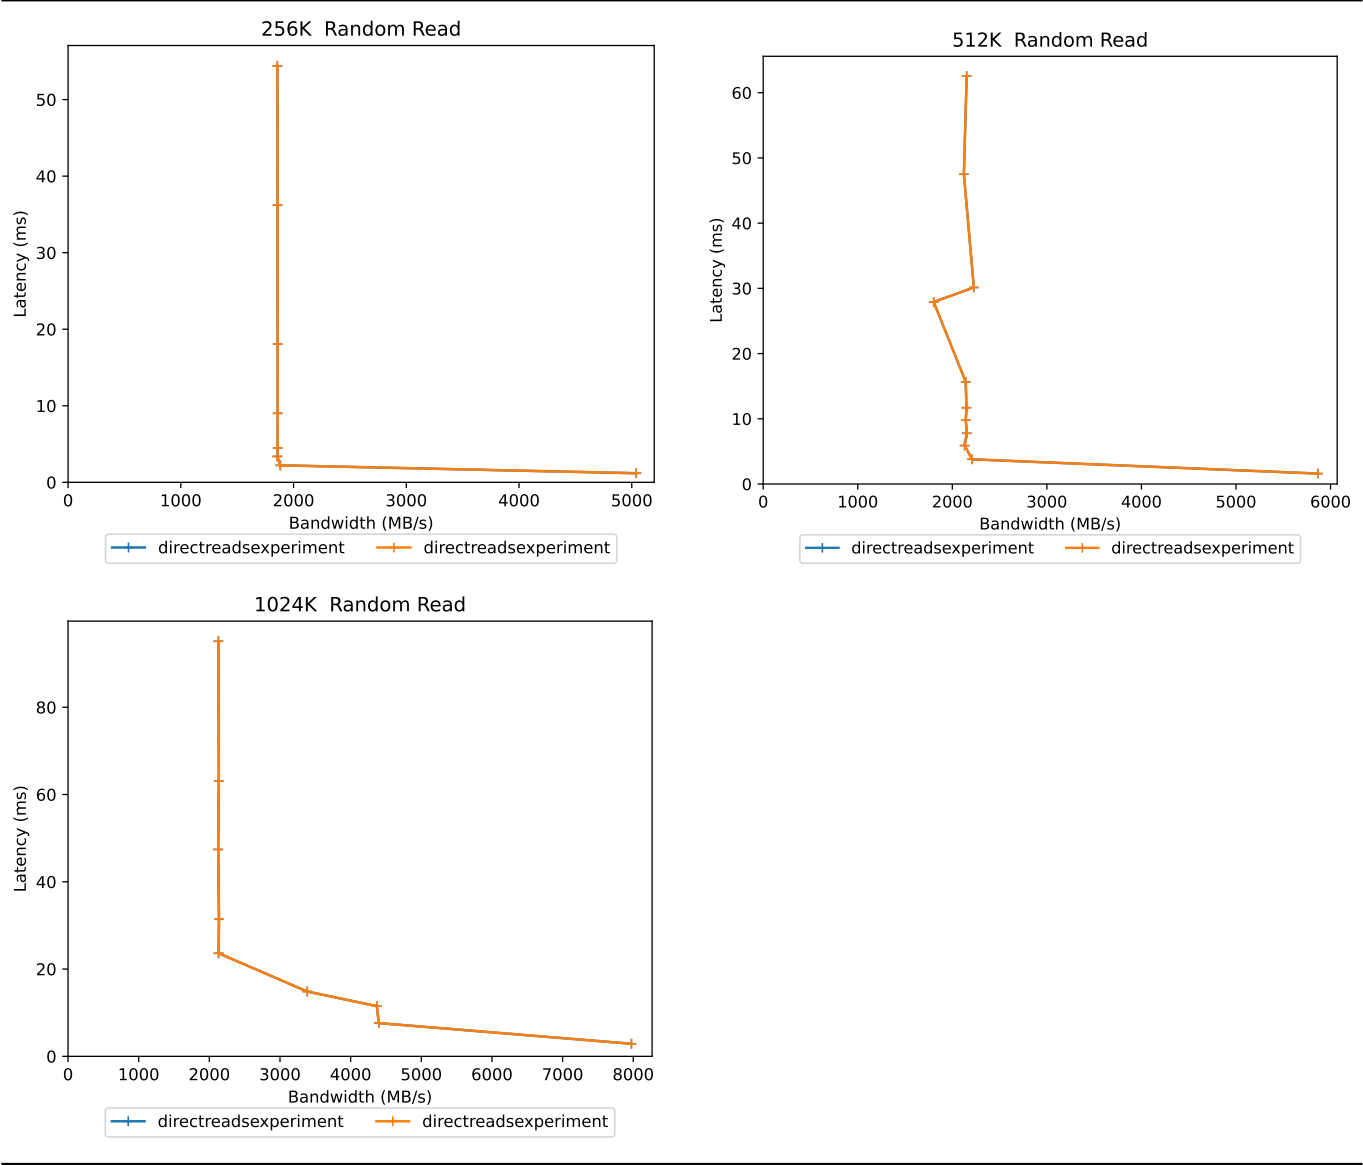
- Comparison summary for directreadsexperiment vs directreadsexperiment
- Response Curves
 - Random Read
- Configuration yaml files
 - results

Comparison summary for directreadsexperiment vs directreadsexperiment

Random Read	directreadsexperiment	directreadsexperiment	%change throughput	%change latency
256K	5040 MB/s@1.2ms	5040 MB/s@1.2ms	0%	0%
512K	5867 MB/s@1.6ms	5867 MB/s@1.6ms	0%	0%
1024K	7974 MB/s@2.9ms	7974 MB/s@2.9ms	0%	0%

Response Curves

Random Read



Configuration yaml files

Only yaml files that differ by more than 20 lines from the yaml file for the baseline directory will be added here in addition to the baseline yaml

results

```

librbd fio:
  cmd_path: /usr/local/bin/fio2
  create_report: true
  fio_out_format: json
  log_avg_msec: 100
  log_bw: true
  log_iops: true
  log_lat: true
  norandommap: true
  osd_ra:
    - 4096
  poolname: rbd_replicated
  prefill:
    blocksize: 64k
    numjobs: 1
  procs_per_volume:
    - 1
  ramp: 20
  rbdname: cbt-librbd fio
  time: 60
  time_based: true
  use_existing_volumes: true
  vol_size: 1000
  volumes_per_client:
    - 8
  wait_pgautoscaler_timeout: 10
  workloads:
    1Mrandomread:
      jobname: randread
      mode: randread
      numjobs:
        - 1
      op_size: 1048576
      total_iodepth:
        - 1
        - 2
        - 3
        - 4
        - 8
        - 12
        - 16
        - 20
        - 24
        - 28
        - 32
        - 48
        - 64
        - 96
        - 128
        - 192
    256krandomread:
      jobname: randread
      mode: randread
      numjobs:
        - 1
      op_size: 262144
      total_iodepth:
        - 1
        - 2
        - 3
        - 4

```

```

- 5
- 8
- 16
- 24
- 32
- 64
- 128
- 256
- 384
512krandomread:
  jobname: randread
  mode: randread
  numjobs:
    - 1
  op_size: 524288
  total_iodepth:
    - 1
    - 2
    - 3
    - 4
    - 8
    - 16
    - 24
    - 32
    - 40
    - 48
    - 64
    - 96
    - 128
    - 192
    - 256
cluster:
  archive_dir: /tmp/cbt
  ceph_mgr_cmd: /usr/bin/ceph-mgr
  ceph_mon_cmd: /usr/bin/ceph-mon
  ceph_osd_cmd: /usr/bin/ceph-osd
  ceph_run_cmd: /usr/bin/ceph-run
  ceph_cmd: /usr/bin/ceph
  clients:
    - --- server1 ---
  clusterid: ceph
  conf_file: /cbt/ceph.conf.4x1x1.fs
  fs: xfs
  head: --- server1 ---
  iterations: 1
  mgrs:
    --- server1 ---:
      a: null
  mkfs_opts: -f -i size=2048
  mons:
    --- server1 ---:
      a: --- IP Address ---:6789
  mount_opts: -o inode64,noatime,logbsize=256k
  osds:
    - --- server1 ---
  osds_per_node: 6
  pdsh_ssh_args: -a -x -l%u %h
  rados_cmd: /usr/bin/rados
  rbd_cmd: /usr/bin/rbd
  tmp_dir: /tmp/cbt
  use_existing: true
  user: root
  monitoring_profiles:

```

```

collectl:
  args: -c 18 -sCD -i 10 -P -oz -F0 --rawtoo --sep ";" -f {collectl_dir}

librbdfio:
  cmd_path: /usr/local/bin/fio2
  create_report: true
  fio_out_format: json
  log_avg_msec: 100
  log_bw: true
  log_iops: true
  log_lat: true
  norandommap: true
  osd_ra:
    - 4096
  poolname: rbd_replicated
  prefill:
    blocksize: 64k
    numjobs: 1
  procs_per_volume:
    - 1
  ramp: 20
  rbdname: cbt-librbdfio
  time: 60
  time_based: true
  use_existing_volumes: true
  vol_size: 1000
  volumes_per_client:
    - 8
  wait_pgautoscaler_timeout: 10
  workloads:
    1Mrandomread:
      jobname: randread
      mode: randread
      numjobs:
        - 1
      op_size: 1048576
      total_iodepth:
        - 1
        - 2
        - 3
        - 4
        - 8
        - 12
        - 16
        - 20
        - 24
        - 28
        - 32
        - 48
        - 64
        - 96
        - 128
        - 192
    256krandomread:
      jobname: randread
      mode: randread
      numjobs:
        - 1
      op_size: 262144
      total_iodepth:
        - 1
        - 2
        - 3

```



```

- 4
- 5
- 8
- 16
- 24
- 32
- 64
- 128
- 256
- 384
512krandomread:
  jobname: randread
  mode: randread
  numjobs:
    - 1
  op_size: 524288
  total_iodepth:
    - 1
    - 2
    - 3
    - 4
    - 8
    - 16
    - 24
    - 32
    - 40
    - 48
    - 64
    - 96
    - 128
    - 192
    - 256
cluster:
  archive_dir: /tmp/cbt
  ceph_mgr_cmd: /usr/bin/ceph-mgr
  ceph_mon_cmd: /usr/bin/ceph-mon
  ceph_osd_cmd: /usr/bin/ceph-osd
  ceph_run_cmd: /usr/bin/ceph-run
  ceph_cmd: /usr/bin/ceph
  clients:
    - --- server1 ---
  clusterid: ceph
  conf_file: /cbt/ceph.conf.4x1x1.fs
  fs: xfs
  head: --- server1 ---
  iterations: 1
  mgrs:
    --- server1 ---:
      a: null
  mkfs_opts: -f -i size=2048
  mons:
    --- server1 ---:
      a: --- IP Address ---:6789
  mount_opts: -o inode64,noatime,logbsize=256k
  osds:
    - --- server1 ---
  osds_per_node: 6
  pdsh_ssh_args: -a -x -l%u %h
  rados_cmd: /usr/bin/rados
  rbd_cmd: /usr/bin/rbd
  tmp_dir: /tmp/cbt
  use_existing: true
  user: root

```

```
monitoring_profiles:
  collect1:
    args: -c 18 -sCD -i 10 -P -oz -F0 --rawtoo --sep ";" -f {collect1_dir}
```